

Lab Outline

Online

This document

You can find the most up-to-date version of this document online [here](#)

(https://github.com/mjbright/jupyter_notebooks/blob/master/2016-Feb_Docker_Build_Lab/2016-Feb_Docker_Build_Lab.md) or as

- a PDF file (https://raw.githubusercontent.com/mjbright/jupyter_notebooks/master/2016-Feb_Docker_Build_Lab/2016-Feb_Docker_Build_Lab.pdf) (<http://bit.ly/1QF0XaH>) or as

- a Jupyter (<http://www.jupyter.org>) notebook at [2016-Feb_Docker_Build_Lab](https://github.com/mjbright/jupyter_notebooks/blob/master/2016-Feb_Docker_Build_Lab/) (https://github.com/mjbright/jupyter_notebooks/blob/master/2016-Feb_Docker_Build_Lab/).

This notebook is runnable in a Jupyter installation with the bash_kernel installed.

Although that is not the subject of this lab, if you want to create your own environment in which to run this lab with Docker components already installed (and even Jupyter/bash_kernel), refer to the README.md [here](https://github.com/mjbright/jupyter_notebooks/blob/master/2016-Feb_Docker_Build_Lab/) (https://github.com/mjbright/jupyter_notebooks/blob/master/2016-Feb_Docker_Build_Lab/)

[in mjbright \(https://www.linkedin.com/in/mjbright\)](https://www.linkedin.com/in/mjbright)
[mjbright \(https://github.com/mjbright\)](https://github.com/mjbright)
[@mjbright \(http://twitter.com/mjbright\)](http://twitter.com/mjbright)

Lab-Description

[TOP](#)

We first need to recuperate the source code examples:

Lab Start

Start this lab by first performing the below step:

In [1]:

```
## Lab Start:

rm -rf ~/src

cd
git clone https://github.com/mjbright/docker-examples src

./src/START_LAB.sh
```

Cloning into 'src'...

remote: Counting objects: 156, done.

remote: Compressing objects: 100% (116/116), done.

remote: Total 156 (delta 64), reused 120 (delta 31), pack-reused 0

Receiving objects: 100% (156/156), 15.60 KiB | 0 bytes/s, done.

Resolving deltas: 100% (64/64), done.

Checking connectivity... done.

Then proceed with the following sections:

- [1. Introduction](#)
- [2. Basic Docker Builds](#)
- [3. Creating Small Images](#)
 - [Creating a small binary with C](#)
 - [Creating a small binary with Go](#)
 - [Creating a toolset Docker image containing several executables](#)
- [4. Pushing our image to Docker Hub](#)
- [5. Dockerfile best practices](#)
- [6. Using the official 'Language Stack' images](#)
 - [Using a Language Stack \(Node.js\)](#)
 - [Using a Language Stack \(Python\)](#)
- [7. Using Compose](#)
 - [Building complex systems with Compose](#)
 - [Rails example with Compose](#)
- [8. Building Docker](#)
 - [Building Docker with Docker](#)

[References](#)

Overall description of the lab steps

NOTE: All lab steps can be considered optional, attendees may perform them in order, or jump to the section of interest to them (to get to the more complicated steps)

Introduction

A refresh on Docker concepts

You may want to skip this section if you have already run the introductory lab.

Look at what docker version you are running. Note that the 'docker version' command reports the local client version as well as the server (docker engine) version.

In [2]:

```
docker version
```

Client:

```
Version:      1.10.0
API version:  1.22
Go version:   go1.5.3
Git commit:   590d5108
Built:       Thu Feb  4 19:55:25 2016
OS/Arch:     linux/amd64
```

Server:

```
Version:      1.10.0
API version:  1.22
Go version:   go1.5.3
Git commit:   590d5108
Built:       Thu Feb  4 19:55:25 2016
OS/Arch:     linux/amd64
```

Images are image layers

Remember that when we talk of a container image it is really a collection of image layers.

The docker info command provides information about the docker engine, see below.

In [4]:

```
docker info
```

```
Containers: 1
  Running: 0
  Paused: 0
  Stopped: 1
Images: 7
Server Version: 1.10.0
Storage Driver: aufs
  Root Dir: /var/lib/docker/aufs
  Backing Filesystem: extfs
  Dirs: 30
  Dirperm1 Supported: false
Execution Driver: native-0.2
Logging Driver: json-file
Plugins:
  Volume: local
  Network: null host bridge
Kernel Version: 3.13.0-77-generic
Operating System: Ubuntu 14.04.3 LTS
OSType: linux
Architecture: x86_64
CPUs: 1
Total Memory: 1.955 GiB
Name: vagrant-ubuntu-trusty-64
ID: 3BDW:ZNNN:SJXZ:MAT3:WBKG:4RQT:PEZJ:7PDN:VG72:XGAL:CU6U:XWX4
WARNING: No swap limit support
```

But if we look at the number of containers and images, the number of images it is not the same as provided above. Why do you think that is?

First let's list the number of running and number of stopped containers

NOTE: the value on your system will be different

In [5]:

```
# Show the running containers:
docker ps

# Count the number of running containers:
echo
echo "Total number of running containers:"
docker ps | tail -n +2 | wc -l
```

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	

```
Total number of running containers:
0
```

In [6]:

```
# Show all the containers (running or stopped):
docker ps -a

# Count all the containers (running or stopped):
echo
echo "Total number of containers (running or stopped):"
docker ps -a | tail -n +2 | wc -l # Number of stopped and running containers ('tail -
```

CONTAINER ID	IMAGE	COMMAND	CREATED
b19891aa2715	lab/c_prog	"/helloWorld"	24 seconds ago
nauseous_mest	Exited (14) 23 seconds ago		

```
Total number of containers (running or stopped):
1
```

We can see that the number of containers reported by docker info correctly reports the number of total containers, running or not

But listing images gives a different value from the 'docker info' value

In [7]:

```
# Show the images:
docker images

# Count the images:
echo
echo "Total number of images:"
docker images | tail -n +2 | wc -l
```

REPOSITORY	TAG	IMAGE ID	CREATED
lab/c_prog	latest	659400ddd8d4	31 seconds ago
swarm	latest	17cf22083736	3 days ago
node	latest	baa18fdeb577	12 days ago
swarm	1.1.0-rc2	81883ac55ffe	2 weeks ago
alpine	latest	14f89d0e6257	2 weeks ago

```
Total number of images:
5
```

That is because there are many intermediate image layers which are not normally listed. But we can list those layers using the '-a' option and now we see a number close to the value from 'docker info'.

(We will see later how the 'docker history' command allows us to see how the layers were created).

In [8]:

```
# Show all the image layers:
docker images -a

# Count all the image layers:
echo
echo "Total number of image layers:"

docker images -a | tail -n +2 | wc -l # The number of image layers+1 (inc. header li
```

REPOSITORY	TAG	IMAGE ID	CREATED
SIZE			
<none>	<none>	8e9605157ccf	42 seconds
ago 0 B			
<none>	<none>	bde06bb6dcf5	42 seconds
ago 877.2 kB			
lab/c_prog	latest	659400ddd8d4	42 seconds
ago 877.2 kB			
swarm	latest	17cf22083736	3 days ago
18.11 MB			
node	latest	baa18fdeb577	12 days ago
643.1 MB			
swarm	1.1.0-rc2	81883ac55ffe	2 weeks ago
18.06 MB			
alpine	latest	14f89d0e6257	2 weeks ago
4.794 MB			

Total number of image layers:
7

Images can include 1 static binary file or more and can even include a whole distribution. Launching a container launches a single process within that container - which may in turn span other child processes.

Let us look at an extremely small image to have an idea just how small an executable image can be. Docker provide an official 'hello-world' image which simply echoes some output to the console.

Let's run that image to see and then investigate the image. First let's search for the image; we see that the first image is 'hello-world' which is an official build

In [9]:

```
docker search hello-world
```

NAME	STARS	OFFICIAL	AUTOMATED	DESCRIPTION
hello-world				Hello World! (an example of mi
nimal Docker...	48		[OK]	
tutum/hello-world				Image to test docker deploymen
ts. Has Apac...	19			[OK]
marcells/aspnet-hello-world				ASP.NET vNext - Hello World
2			[OK]	
bonomat/nodejs-hello-world				a simple nodejs hello world co
ntainer	1			[OK]
carinamarina/hello-world-app				This is a sample Python web ap
plication, r...	1			[OK]
vegasbrianc/docker-hello-world				
1			[OK]	
carinamarina/hello-world-web				A Python web app, running on p
ort 5000, wh...	1			[OK]
wowgroup/hello-world				Minimal web app for testing pu
rposes	0			[OK]
dockercloud/hello-world				Sample hello world application
0			[OK]	
crccheck/hello-world				Hello World web server in unde
r 2.5 MB	0			[OK]
bencampbell/hello-world				First automated build.
0			[OK]	
nirmata/hello-world				
0			[OK]	
n8io/hello-world				A simple hello world node.js a
pp to test d...	0			[OK]
mikelh/hello-world				simplified hello world as dumm
y start for ...	0			[OK]
poojathote/hello-world				this is 3rd POC
0			[OK]	
vamosporpartes/hello-world				Primeiro projeto para testar o
link entre ...	0			[OK]
asakaguchi/docker-nodejs-hello-world				Hello World for Docker
0			[OK]	
cpro/http-hello-world				Hello world
0			[OK]	
ileontyev81/docker-hello-world				hello world test build
0			[OK]	
chaliceg/docker-hello-world				Hello world
0			[OK]	
alexwelch/hello-world				
0			[OK]	
vasia/docker-hello-world				rhrrthrh
0			[OK]	
asakaguchi/magellan-nodejs-hello-world				Hello World for MAGELLAN
0			[OK]	
wodge/docker-hello-world				Hello World test for auto upda
te to Docker...	0			[OK]
chalitac/hello-world				Just Hello World
0			[OK]	

Let's now run that image

In [10]:

```
docker run hello-world
```

Note how we see the pulling of the image if not already available locally:

```
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
```

```
03f4658f8b78: Pulling fs layer
```

```
Digest: sha256:8be990ef2aeb16dbcb9271ddfe2610fa6658d13f6dfb8bc72074cc1c
a36966a7
```

```
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker.

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

<https://hub.docker.com> (<https://hub.docker.com>)

For more examples and ideas, visit:

<https://docs.docker.com/userguide/> (<https://docs.docker.com/userguide/>)

If it took a while to run, this was due to the time needed to download the image before running it - see above.

Try the command a second time to see how it runs instantaneously as there is no need to download the image which already exists locally on the 'docker engine'.

In [11]:

```
docker run hello-world
```

```
# The second time there is no need to repull the image:
```

Hello from Docker.

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

<https://hub.docker.com> (<https://hub.docker.com>)

For more examples and ideas, visit:

<https://docs.docker.com/userguide/> (<https://docs.docker.com/userguide/>)

Let us inspect the image. We see that the file is only 960 bytes large, it must be machine code to print out the text. So we see that an image can be really very small

In [12]:

```
docker images hello-world
```

REPOSITORY	TAG	IMAGE ID	CREATED
hello-world	latest	690ed74de00f	3 months ago
o	960 B		

We can also inspect the image with the history command to see how it was constructed.

Note that history shows the image layers in reverse order, latest first.

From the below command we can see that the image was created from only 2 image layers.

The image was built simply by copying in a binary executable and then specifying the default command to invoke when the image is run.

In [13]:

```
docker history hello-world
```

IMAGE SIZE	CREATED COMMENT	CREATED BY
690ed74de00f	3 months ago	/bin/sh -c #(nop) CMD ["/hell o"]
<missing>	3 months ago	/bin/sh -c #(nop) COPY file:1ad 52e3eaf4327c8f 960 B

In [14]:

```
echo
echo "Total size (in bytes) of text in 'hello-world' image:"
docker run hello-world | wc -c
```

Total size (in bytes) of text in 'hello-world' image:
801

So we see that 801 bytes of that executable is the actual text printed ! So the real program size is roughly 160 bytes (of assembler no doubt)

Basic Docker Builds

[TOP](#)

Dockerfile

Images are built from Dockerfiles which contain a series of commands used to build up a docker image. Note that each command in the Dockerfile results in a new image layer being created, no matter how trivial the command - even ENV "commands" create a new image layer.

In the following lab we will see how images can be built systematically from a Dockerfile using the 'docker build' command.

DockerHub

When we pull an image we pull it from a Docker Registry. The [DockerHub \(https://hub.docker.com/\)](https://hub.docker.com/) is a free to use Docker registry allowing to store your own image files (which are publicly available unless you pay for your account) and to pull other image files of other users or officially provided images.

You can create images either by

- building them from a Dockerfile (thus in a **repeatable** manner)
- building them manually by modifying a running container and '*commit*'ing it's state

The DockerHub contains images which may be

- **Automated builds** (built from a git repository)
 - Such builds are usually built from an open-source git repo and so are called **Trusted builds** because the source code is available. *Note:* The github repo may contain binary files though
- **Official builds** are builds which are provided by partners or by Docker themselves

Other images may exist in the hub but their origin is unknown and so represent a security risk.

It is possible to search the DockerHub, or another Docker Registry, using the 'docker search' command with appropriate options. Other companies offer their own Docker Registry which may be freely accessible e.g. RedHat, internal to a company e.g. HPE IT, or available as part of a paid for service e.g. IBM or Amazon Web Services ECS.

In [20]:

```
cd ~/src/basic
```

In the ~/test folder create a Dockerfile with the contents shown below (the o/p of the cat command).

For this you may use vi if you are familiar, otherwise the 'nano' text editor is recommended.

Use ctrl-W to write out the file and ctrl-X to quit the editor.

In [25]:

```
cat Dockerfile
```

```
#
# Dockerfile to demonstrate the simplest build
#

FROM python

MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

# NOTE: all RUN commands are executed at build time, look at the output
# of the "docker build"
# below and you will see the output of the following line when we
# build the first time
# (i.e. when caching is not used):

RUN echo DATE: This build is being done at $(date)
#RUN python --version

CMD bash
```

In [26]:

```
ls -altr Dockerfile
```

```
-rw-rw-r-- 1 vagrant vagrant 439 Feb  8 15:31 Dockerfile
```

We can now build a new image using this dockerfile using the below command where

- we specify the current directory as the context for the build (any ADD/COPY or Dockerfile files will be sourced from here) with the '.' option
- we specify the specific tag to use for the generated image as *"lab/basic"* with

-t lab/basic

In [27]:

```
docker build -t lab/basic .
```

Sending build context to Docker daemon 2.048 kB

Step 1 : FROM python

---> 93049cc049a6

Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

---> Running in 1e836ce0bd8b

---> 265f1e4a525d

Removing intermediate container 1e836ce0bd8b

Step 3 : RUN echo DATE: This build is being done at \$(date)

---> Running in e195f1a8d6fc

DATE: This build is being done at Mon Feb 8 15:31:23 UTC 2016

---> 58100f46a8de

Removing intermediate container e195f1a8d6fc

Step 4 : CMD bash

---> Running in 6a24b4f9634c

---> c028a15ddd50

Removing intermediate container 6a24b4f9634c

Successfully built c028a15ddd50

Note that during the build, the RUN commands are actually run.

They are used to build up this new image.

In this case we echo the 'Python' version string during the build process.

You can see the available options to the build command by issuing 'docker build --help'

In [28]:

```
docker build --help
```

Usage: docker build [OPTIONS] PATH | URL | -

Build an image from a Dockerfile

<code>--build-arg=[]</code>	Set build-time variables
<code>--cpu-shares</code>	CPU shares (relative weight)
<code>--cgroup-parent</code>	Optional parent cgroup for the container
<code>--cpu-period</code>	Limit the CPU CFS (Completely Fair Scheduler) period
<code>--cpu-quota</code>	Limit the CPU CFS (Completely Fair Scheduler) quota
<code>--cpuset-cpus</code>	CPU(s) in which to allow execution (0-3, 0,1)
<code>--cpuset-mems</code>	MEM(s) in which to allow execution (0-3, 0,1)
<code>--disable-content-trust=true</code>	Skip image verification
<code>-f, --file</code>	Name of the Dockerfile (Default is 'PATH/Dockerfile')
<code>--force-rm</code>	Always remove intermediate containers
<code>--help</code>	Print usage
<code>--isolation</code>	Container isolation level
<code>-m, --memory</code>	Memory limit
<code>--memory-swap</code>	Swap limit equal to memory plus swap: '-1' to enable unlimited swap
<code>--no-cache</code>	Do not use cache when building the image
<code>--pull</code>	Always attempt to pull a newer version of the image
<code>-q, --quiet</code>	Suppress the build output and print image ID on success
<code>--rm=true</code>	Remove intermediate containers after a successful build
<code>--shm-size</code>	Size of /dev/shm, default value is 64 MB
<code>-t, --tag=[]</code>	Name and optionally a tag in the 'name:tag' format
<code>--ulimit=[]</code>	Ulimit options

We can see all the images available using the 'docker images' command

but if there are many, how do we see just our newly-created image?

You can see the available options to the images command by issuing 'docker images --help'

In [29]:

```
docker images --help
```

Usage: docker images [OPTIONS] [REPOSITORY[:TAG]]

List images

-a, --all	Show all images (default hides intermediate images)
--digests	Show digests
-f, --filter=[]	Filter output based on conditions provided
--format	Pretty-print images using a Go template
--help	Print usage
--no-trunc	Don't truncate output
-q, --quiet	Only show numeric IDs

So you can see your newly built 'lab/basic' with the following command:

In [30]:

```
docker images lab/basic
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
lab/basic 689.1 MB	latest	c028a15ddd50	18 seconds

Note that if you rerun the build command, the build should run faster, you will notice how build steps recognize that this step has already been performed and so will use the image layer already available in the local cache.

Now let us see what happens if we modify our Dockerfile, by inserting a line, such as defining an environment variable.

We will use the same Dockerfile, but this time we will insert an "ENV" line

In [31]:

```
cd ~/src/basic/
```

Now edit the Dockerfile to have the contents as shown below (the o/p of the cat command).

In [32]:

```
cat Dockerfile
```

```
#
# Dockerfile to demonstrate the simplest build
#

FROM python

MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

# NOTE: all RUN commands are executed at build time, look at the output
of the "docker build"
#     below and you will see the output of the following line when we
build the first time
#     (i.e. when caching is not used):

RUN echo DATE: This build is being done at $(date)
#RUN python --version

ENV myVar=anythingReally

CMD bash
```

This time when we build the image we will see that the addition of a line between the "RUN" line and the "CMD" line forces rebuild of subsequent image layers.

We see 'Using cache' for Step 2 and 3 only

In [33]:

```
docker build -t lab/basic .
```

```
Sending build context to Docker daemon 2.048 kB
Step 1 : FROM python
---> 93049cc049a6
Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>
---> Using cache
---> 265f1e4a525d
Step 3 : RUN echo DATE: This build is being done at $(date)
---> Using cache
---> 58100f46a8de
Step 4 : ENV myVar anythingReally
---> Running in 86c9262af943
---> f493d0962635
Removing intermediate container 86c9262af943
Step 5 : CMD bash
---> Running in 42344c05af2a
---> 6566cf6ed29d
Removing intermediate container 42344c05af2a
Successfully built 6566cf6ed29d
```

Similarly we can force to not use the cache with the --no-cache option.

This could be useful if we suspect the caching is not working properly due to some external change.

In [34]:

```
docker build --no-cache -t lab/basic .
```

Sending build context to Docker daemon 2.048 kB

Step 1 : FROM python

---> 93049cc049a6

Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

---> Running in 352ab6578552

---> c5efd78c5c14

Removing intermediate container 352ab6578552

Step 3 : RUN echo DATE: This build is being done at \$(date)

---> Running in 2be6e59360a4

DATE: This build is being done at Mon Feb 8 15:35:15 UTC 2016

---> 77f9a84551cc

Removing intermediate container 2be6e59360a4

Step 4 : ENV myVar anythingReally

---> Running in f56c83b6c061

---> e7e8c2a26251

Removing intermediate container f56c83b6c061

Step 5 : CMD bash

---> Running in afdea0bc90b9

---> e3d7946e9d93

Removing intermediate container afdea0bc90b9

Successfully built e3d7946e9d93

In [35]:

```
docker images lab/basic
```

REPOSITORY	TAG	IMAGE ID	CREATED
lab/basic	latest	e3d7946e9d93	12 seconds
ago	689.1 MB		

Creating small images

[TOP](#)

Creating a small C Docker image

[TOP](#)

In this example we show how we can create a Docker image from a statically-linked binary.

The goal of this step is to show that we do not need an Operating System image for a Docker container.

All we need is a self-contained binary - i.e. statically linked binary.

Of course a dynamically linked binary could also be used, but in this case it's more complicated as you would have to manually add all it's dependent libraries. Let's let gcc to do that work for us!

This section comprises 2 things

- A Dockerfile to build our image from a static binary Note that it starts with "FROM scratch".
Scratch is a special 'empty' image
- helloFromDocker.c

So first let's build our static binary

In [36]:

```
cd ~/src/createTinyC/  
  
# For RHEL/Fedora/Centos only:  
# First we must install *glibc-static*  
#yum install -y glibc-static  
  
gcc -static helloWorld.c -o helloWorld  
  
ls -alh helloWorld
```

```
-rwxrwxr-x 1 vagrant vagrant 857K Feb  8 15:35 helloWorld
```

So we see that this created a binary file of approximately 857kby.

Now let's build our Docker image containing this binary.

You will need to recreate the Dockerfile as follows:

In [37]:

```
cat Dockerfile
```

```
FROM scratch  
MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>  
  
ADD ./helloWorld /helloWorld  
CMD ["/helloWorld"]
```

In [38]:

```
docker build -t lab/c_prog .
```

Sending build context to Docker daemon 882.2 kB

Step 1 : FROM scratch

--->

Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

---> Using cache

---> 8e9605157ccf

Step 3 : ADD ./helloWorld /helloWorld

---> Using cache

---> bde06bb6dcf5

Step 4 : CMD /helloWorld

---> Using cache

---> 659400ddd8d4

Successfully built 659400ddd8d4

If we now look at the generated Docker image (below) we see an image of about 877kby.

So whilst this is larger than the 1kby hello-world image (no doubt written in assembler) it is still a very small Docker image which is only 20kbytes larger than the original binary file.

In [39]:

```
docker images lab/c_prog
```

REPOSITORY	TAG	IMAGE ID	CREATED
lab/c_prog	latest	659400ddd8d4	17 hours ag
o	877.2 kB		

And now let's run that image

In [40]:

```
docker run lab/c_prog
```

Hello World!!

In [41]:

```
docker history lab/c_prog
```

IMAGE	CREATED	CREATED BY
659400ddd8d4	17 hours ago	/bin/sh -c #(nop) CMD ["/hellow
orld"]	0 B	
bde06bb6dcf5	17 hours ago	/bin/sh -c #(nop) ADD file:eb1f
2e0a74291359d0	877.2 kB	
8e9605157ccf	17 hours ago	/bin/sh -c #(nop) MAINTAINER "D
ocker Build La	0 B	

Creating a small Go Docker image

[TOP](#)

That's fine, but isn't Go taking over the world as a systems language? Docker, Kubernetes, LXD, Rocket, ... many new tools are being written in Go.

Let's see how we can do the same exercise but building a Go statically-linked binary.

The goal of this step is as the previous step (building an image from a single statically-linked binary) but using Go, but also to demonstrate how we can use a Docker image containing a Go compiler, rather than explicitly installing a compiler.

NOTE: We will do this **without** 'installing a Go compiler'

In [42]:

```
cd ~/src/createTinyGo
cat Dockerfile
```

```
FROM scratch
MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>
```

```
ADD ./hello /hello
CMD ["/hello"]
```

NOW we invoke the golang container to build our go source code.

The following docker run

- mounts the current directory (\$PWD) as /go within the container
- launches a container of the **golang** image which contains the go compiler
- invokes the command "go build -v hello" on the container to build the sources for the "hello.go" code.

The hello.go code is located under src/hello/hello.go.

This is a Go convention.

NOTE: The important thing to note here is that the compiler is within the image. We did not need to install a native Go compiler, we used an image which contains the compiler and by mounting the current directory the container can read the source code and write the executable outside the container. This is a nice pattern of providing a tool within a container.

In [43]:

```
# Compile our Go-code using the Go compiler provided by the 'golang' container:
docker run -it -v $PWD:/go golang go build hello

# Now we have our small static binary compiled
ls -l hello
```

Unable to find image 'golang:latest' locally
latest: Pulling from library/golang

```
03e1855d4f31: Already exists
a3ed95caeb02: Already exists
9269ba3950bb: Already exists
6ecee6444751: Already exists
a3ed95caeb02: Pulling fs layer
8ccc76a74787: Pulling fs layer
70fd525fbd3f: Pulling fs layer
Digest: sha256:2242a664699c7dd1aaf3c380107d20bfff45cbbe39d92e848c4c83240
adc4169f
Status: Downloaded newer image for golang:latest
-rwxr-xr-x 1 root root 2367272 Feb  8 15:37 hello
```

Now we can build our image including this static binary.

In [44]:

```
docker build -t lab/go-hello .
```

```
Sending build context to Docker daemon 2.377 MB
Step 1 : FROM scratch
--->
Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>
---> Using cache
---> 8e9605157ccf
Step 3 : ADD ./hello /hello
---> 0aac2ea7bdeb
Removing intermediate container e2ace3cfae0c
Step 4 : CMD /hello
---> Running in 345e7db6224d
---> 438d3567ae8a
Removing intermediate container 345e7db6224d
Successfully built 438d3567ae8a
```

In [45]:

```
docker images lab/*
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
lab/go-hello go 2.367 MB	latest	438d3567ae8a	4 seconds a
lab/basic go 689.1 MB	latest	e3d7946e9d93	2 minutes a
lab/c_prog o 877.2 kB	latest	659400ddd8d4	17 hours ag

Creating a toolset Docker image containing several executables

[TOP](#)

Now let's see how we can combine these static binaries into one image.

Let's build a new image derived from the Docker provided 'hello-world' image

The goal of this step is to show how we can combine several executables in an image, opening up the possibility of creating a container of tools.

We will do this without directly 'installing a Go compiler' but by using the official '*golang*' image which includes the Go compiler.

In [46]:

```
cd ~/src/toolset

cp ../createTinyC/helloWorld    helloWorld
cp ../createTinyGo/hello        helloWorldGo

ls -altr
```

```
total 3192
-rw-rw-r-- 1 vagrant vagrant    68 Feb  7 21:27 helloWorld.c
-rw-rw-r-- 1 vagrant vagrant   181 Feb  7 21:32 Dockerfile
-rwxrwxr-x 1 vagrant vagrant   333 Feb  7 21:32 createTinyDockerImag
e.sh
drwxrwxr-x 11 vagrant vagrant  4096 Feb  8 15:02 ..
-rwxrwxr-x 1 vagrant vagrant 877192 Feb  8 15:38 helloWorld
drwxrwxr-x 2 vagrant vagrant   4096 Feb  8 15:38 .
-rwxr-xr-x 1 vagrant vagrant 2367272 Feb  8 15:38 helloWorldGo
```

Create the Dockerfile with the following contents

In [47]:

```
cat Dockerfile
```

```
FROM hello-world
MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

ADD ./helloWorld /helloWorld
CMD ["/helloWorld"]

ADD ./helloWorldGo /helloWorldGo
CMD ["/helloWorldGo"]
```

Now build the toolset with these executables included:

In [48]:

```
docker build -t lab/toolset ./
```

```
Sending build context to Docker daemon  3.25 MB
Step 1 : FROM hello-world
---> 690ed74de00f
Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>
---> Running in 2a30569f66fa
---> 2813581e3ac0
Removing intermediate container 2a30569f66fa
Step 3 : ADD ./helloWorld /helloWorld
---> e8d91e3e5713
Removing intermediate container 0a36f6762299
Step 4 : CMD /helloWorld
---> Running in 8eba1b0029d1
---> db82d13ae941
Removing intermediate container 8eba1b0029d1
Step 5 : ADD ./helloWorldGo /helloWorldGo
---> 1c25395172b0
Removing intermediate container aa5e10abd507
Step 6 : CMD /helloWorldGo
---> Running in 1ed9542b8551
---> 6b82de537bd1
Removing intermediate container 1ed9542b8551
Successfully built 6b82de537bd1
```

If we look at the history of this image we can see the different executables and CMDs which have been added including the original hello-world image.

In [49]:

```
docker history lab/toolset
```

IMAGE SIZE	CREATED COMMENT	CREATED BY
6b82de537bd1	5 seconds ago	/bin/sh -c #(nop) CMD ["/hello
or1dGo"]	0 B	
1c25395172b0	5 seconds ago	/bin/sh -c #(nop) ADD file:e827
5a9025432fdf2e	2.367 MB	
db82d13ae941	6 seconds ago	/bin/sh -c #(nop) CMD ["/hello
or1d"]	0 B	
e8d91e3e5713	6 seconds ago	/bin/sh -c #(nop) ADD file:eb1f
2e0a74291359d0	877.2 kB	
2813581e3ac0	6 seconds ago	/bin/sh -c #(nop) MAINTAINER "D
ocker Build La	0 B	
690ed74de00f	3 months ago	/bin/sh -c #(nop) CMD ["/hell
o"]	0 B	
<missing>	3 months ago	/bin/sh -c #(nop) COPY file:1ad
52e3eaf4327c8f	960 B	

Now we are free to specify which command is to be run.

If we don't specify the command, the last (first in the above history list) will be run (so /helloWorldGo in this case)

In [50]:

```
docker run lab/toolset
```

```
Hello world from Go !!
```

Or we can explicitly choose the executable to be run, such as the /hello executable of the original "hello-world" image

In [51]:

```
docker run lab/toolset /hello
```

Hello from Docker.

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

<https://hub.docker.com> (<https://hub.docker.com>)

For more examples and ideas, visit:

<https://docs.docker.com/userguide/> (<https://docs.docker.com/userguide/>)

In [52]:

```
docker run lab/toolset /helloWorld
```

Hello World!!

In [53]:

```
docker run lab/toolset /helloWorldGo
```

Hello world from Go !!

We have seen how we can combine several executables in an image, and we can imagine creating a toolset container in this way (with some more useful executable tools!)

Pushing our image to Docker Hub

[TOP](#)

Note: If you have your own account on Docker Hub you may wish to use that for this exercise.

Otherwise we will all be using the same account '**dockerlabs**' so you will need to specify a tag which distinguishes your images from your neighbours.

The goal of this step is to demonstrate how we may push an image which we have built to the Docker Hub.

First we will retag our local image to be unique. If you are on **podN**, then tag with **userN**,

e.g. if you are **pod3**,

```
docker tag lab/toolset dockerlabs/toolset:user3
```

Notice that we then have 2 toolset images with different tags.

They are otherwise identical (but they could be different) and have the same "IMAGE ID".

In [54]:

```
docker tag lab/toolset:latest dockerlabs/toolset:userN
docker images */toolset
```

REPOSITORY	TAG	IMAGE ID	CREATED
SIZE			
dockerlabs/toolset	userN	6b82de537bd1	36 minutes
ago 3.245 MB			
lab/toolset	latest	6b82de537bd1	36 minutes
ago 3.245 MB			

First we must login to the Docker Hub.

Ask you instructor for the password to the dockerlabs account.

In [56]:

```
docker login -u dockerlabs -p $PASSWORD -e dockerlabs@mjbright.net
```

WARNING: login credentials saved in /home/vagrant/.docker/config.json
Login Succeeded

Now we may push our image to the public Docker Hub

In [57]:

```
docker push dockerlabs/toolset:userN
```

The push refers to a repository [docker.io/dockerlabs/toolset]

```
c04609a4ced5: Preparing
9cd6282bb409: Preparing
5f70bf18a086: Preparing
userN: digest: sha256:b77d291ac6e7b86640922f63048dd55f654cc1b2dc8019647
cfbaed20b643cc8 size: 4619
```

NOTE: The docker search command is not very useful.

and the below command doesn't show us the tags ... and so we don't know if the below image is tagged user1, user2, ...

In [58]:

```
docker search dockerlabs/
```

NAME	DESCRIPTION	STARS	OFFICIAL	AUTOMATED
dockerlabs/toolset		0		

Logging on to DockerHub to see your tagged image there

So for this step, log onto DockerHub <https://hub.docker.com/> (<https://hub.docker.com/>)

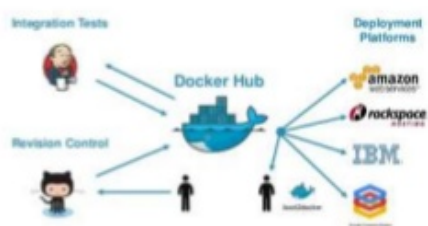
In [59]:

```
# Ignore this line: it is just to display the image below
```

```
curl -s 'http://image.slidesharecdn.com/dockerdemystifiedforsbjug-150918181554-lva1-a
```

Docker Hub

Storage for Docker Containers
Maintains Lineage / All Versions
Public, Private & Self-Hosted Repositories
Like GitHub, but for Docker Images



As dockerlabs (dockerlabs AT mjbright.net) with the appropriate password (ask your instructor)

Once logged you should see the dockerlabs/toolset listed, otherwise you can search for it.

Click on the [dockerlabs/toolset](https://hub.docker.com/r/dockerlabs/toolset/) (<https://hub.docker.com/r/dockerlabs/toolset/>) link, then on the [Tags](https://hub.docker.com/r/dockerlabs/toolset/tags/) (<https://hub.docker.com/r/dockerlabs/toolset/tags/>) link and you should now see your tagged image there.

Remove any running 'dockerlabs/toolset' containers on your system

We do this step to make sure we can easily delete your local dockerlabs/toolset:userN image.

These steps could be done by hand through use of 'docker ps' and 'docker ps -a' and picking containers ids corresponding to 'dockerlabs/toolset' containers to use with 'docker stop' and 'docker rm' commands.

The below expressions do this automatically for us.

In [60]:

```
IMAGE_NAME=dockerlabs/toolset

echo; echo "Currently running or stopped '$IMAGE_NAME' containers"
docker ps -a --filter=ancestor=$IMAGE_NAME

echo; echo "Stopping any running '$IMAGE_NAME' containers (so we can remove dockerlab
docker stop $(docker ps --filter=ancestor=$IMAGE_NAME) 2>/dev/null

echo; echo "Removing any stopped '$IMAGE_NAME' containers (so we can remove dockerlab
docker rm $(docker ps -a --filter=ancestor=$IMAGE_NAME) 2>/dev/null

echo; echo "There should be no more '$IMAGE_NAME' containers present:"
docker ps -a --filter=ancestor=$IMAGE_NAME
```

Currently running or stopped 'dockerlabs/toolset' containers

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	

Stopping any running 'dockerlabs/toolset' containers (so we can remove dockerlabs/ image)

Removing any stopped 'dockerlabs/toolset' containers (so we can remove dockerlabs/ image)

There should be no more 'dockerlabs/toolset' containers present:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	

In [61]:

```
docker images dockerlabs/*
```

REPOSITORY	TAG	IMAGE ID	CREATED
SIZE			
dockerlabs/toolset	userN	6b82de537bd1	38 minutes
ago	3.245 MB		

Note that the following rmi command 'Untags' the image.

This is because it is the same - has the same image id - as our original 'lab/toolset' image.

Removing the dockerlabs/toolset image does not remove the identical 'lab/toolset' image but removes the 'dockerlabs/toolset' tag.

In [62]:

```
docker rmi dockerlabs/toolset:userN
```

Untagged: dockerlabs/toolset:userN

In [63]:

```
docker images dockerlabs/*
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
--------------------	-----	----------	---------

As we have removed ('untagged') the dockerlabs/toolset image, the following run command will download it from the Docker Hub

In [64]:

```
docker run dockerlabs/toolset:userN
```

Unable to find image 'dockerlabs/toolset:userN' locally
userN: Pulling from dockerlabs/toolset

03f4658f8b78: Already exists

a3ed95caeb02: Already exists

157cafd668ce: Already exists

Digest: sha256:b77d291ac6e7b86640922f63048dd55f654cc1b2dc8019647cfbaed2
0b643cc8

Status: Downloaded newer image for dockerlabs/toolset:userN

Hello world from Go !!

In [65]:

```
docker images dockerlabs/*
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
dockerlabs/toolset ago 3.245 MB	userN	d34fbee9ad6f	39 minutes

In [66]:

```
docker run dockerlabs/toolset:userN /helloWorld
```

Hello World!!

In [67]:

```
docker run dockerlabs/toolset:userN /hello
```

Hello from Docker.

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker Hub account:

<https://hub.docker.com> (<https://hub.docker.com>)

For more examples and ideas, visit:

<https://docs.docker.com/userguide/> (<https://docs.docker.com/userguide/>)

Dockerfile best practices

[TOP](#)

The goal of this step is to demonstrate certain Dockerfile optimizations.

- group related commands together using '&&' to reduce image layers
- if temporary files are to be removed

In [68]:

```
cd ~/src/build-best-practices
cat Dockerfile
```

```
FROM ubuntu
```

```
MAINTAINER "Docker Labs" <dockerlabs@mjbright.net>
```

```
#
# Instead of performing the following commands individually which
# involves creating a separate image layer for each RUN command:
#   RUN apt-get update
#   RUN apt-get -y -q upgrade
#   RUN rm -rf /var/lib/apt/lists/*
#
# Here we combine the update, upgrade and cleanup steps into one command
# - This produces less image layers (better for disk space and performance)
# - This keeps image smaller by removing temporary files in the same layer
#   If we performed update/upgrade and then rm as a separate step there would
#   be an intermediate layer including those files, making the overall image larger.
#
RUN apt-get update && apt-get -y -q upgrade && rm -rf /var/lib/apt/lists/*
```

TO be completed ... !!

Using the official 'Language Stack' images

[TOP](#)

Creating a Node.js application from the Node.js 'LanguageStack' Docker image

[TOP](#)

Docker provide a set of '*Language Stacks*' which are medium sized images representing the necessary dependencies for a particular language.

The goal of this step is to demonstrate the use of Docker-provided *Language Stacks*.

On the [Docker Hub \(https://hub.docker.com/\)](https://hub.docker.com/) we can find language stacks available for a variety of languages/environments, each with different release versions (Python 2.x and Python 3.x for example):

- [Node.js \(Javascript\) \(https://hub.docker.com/_/node/\)](https://hub.docker.com/_/node/)
- [Python \(https://hub.docker.com/_/python/\)](https://hub.docker.com/_/python/)
- [Ruby \(https://hub.docker.com/_/ruby/\)](https://hub.docker.com/_/ruby/)

You can browse the complete list of 'Official Images' on the Docker Hub [here \(https://hub.docker.com/explore/\)](https://hub.docker.com/explore/)

Now let's look at an example of Node.js. To run a Node.js application this time we will need

In [69]:

```
docker pull node
```

Using default tag: latest

latest: Pulling from library/node

03e1855d4f31: Already exists

a3ed95caeb02: Already exists

9269ba3950bb: Already exists

6ecee6444751: Already exists

7a0c192d4d25: Already exists

a3ed95caeb02: Already exists

Digest: sha256:1bdda7cdd0a8f9c44ac6f51c77de9f42ed3f62efdf557dba6bcca675084de1bd

Status: Image is up to date for node:latest

In [70]:

```
docker images node
```

REPOSITORY	TAG	IMAGE ID	CREATED
node	latest	baa18fdeb577	12 days ago
643.1 MB			

In [71]:

```
docker history node
```

IMAGE SIZE	CREATED COMMENT	CREATED BY
baa18fdeb577 0 B	12 days ago	/bin/sh -c #(nop) CMD ["node"]
<missing>	12 days ago	/bin/sh -c curl -SLO "https://n
odejs.org/dist	36.39 MB	
<missing>	12 days ago	/bin/sh -c #(nop) ENV NODE_VERS
ION=5.5.0	0 B	
<missing>	12 days ago	/bin/sh -c #(nop) ENV NPM_CONFI
G_LOGLEVEL=inf	0 B	
<missing>	12 days ago	/bin/sh -c set -ex && for key
in 9554F0	51.75 kB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	314.7 MB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	122.6 MB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	44.3 MB	
<missing>	13 days ago	/bin/sh -c #(nop) CMD ["/bin/ba
sh"]	0 B	
<missing>	13 days ago	/bin/sh -c #(nop) ADD file:e5a3
d20748c5d3dd5f	125.1 MB	

In [72]:

```
cd ~/src/nodeJS/
ls -altr
```

```
total 24
drwxrwxr-x 2 vagrant vagrant 4096 Feb 7 21:27 src
-rw-rw-r-- 1 vagrant vagrant 116 Feb 7 21:27 README.md
-rw-rw-r-- 1 vagrant vagrant 315 Feb 7 21:32 Dockerfile
-rwxrwxr-x 1 vagrant vagrant 78 Feb 7 21:32 build_run.sh
drwxrwxr-x 3 vagrant vagrant 4096 Feb 7 21:32 .
drwxrwxr-x 11 vagrant vagrant 4096 Feb 8 15:02 ..
```

Once again edit the Dockerfile to have the contents shown below:

In [73]:

```
cat Dockerfile
```

```
FROM node
```

```
# make the src folder available in the docker image  
ADD src/ /src
```

```
WORKDIR /src
```

```
# install the dependencies from the package.json file  
RUN npm install
```

```
# make port 80 available outside of the image  
EXPOSE 80
```

```
# start node with the index.js file of our hello-world application  
CMD ["node", "index.js"]
```

Now let's build the image

In [74]:

```
docker build -t node-hello .
```

Sending build context to Docker daemon 6.656 kB

Step 1 : FROM node

---> baa18fdeb577

Step 2 : ADD src/ /src

---> 185d94954037

Removing intermediate container 9bb58bd74883

Step 3 : WORKDIR /src

---> Running in ff9be9297185

---> b9d929ce3d90

Removing intermediate container ff9be9297185

Step 4 : RUN npm install

---> Running in 40468215f688

npm info it worked if it ends with ok

npm info using npm@3.3.12

npm info using node@v5.5.0

npm info attempt registry request try #1 at 4:18:42 PM

npm http request GET <https://registry.npmjs.org/express> (<https://registry.npmjs.org/express>)

npm http 200 <https://registry.npmjs.org/express> (<https://registry.npmjs.org/express>)

and run the image in the background, exposing port 80

In [75]:

```
docker run -p 80:80 --name web -d node-hello
```

c43f44797eddb266ffa9d9d69fa91e12927284f8ec0af209193865ca3a43b434

Now let's use curl to access this container (default port for curl is 80)

In [76]:

```
curl http://localhost
```

<html><body>Hello from Node.js container c43f44797edd</body></html>

Creating a Python application from the Python 'LanguageStack' Docker image

[TOP](#)

The goal of this step is to demonstrate the use of the Python *Language Stack*.

Now let's look at a Python example. To run a Node.js application this time we will need

Let's pull and examine the official 'Docker Language Stack' image of Python

Note how the earliest image layers (at the bottom of the list) have the same image ids as the earliest image layers of the Node.js image.

So we can see that they were both created from the same base.

In [77]:

```
docker pull python
```

Using default tag: latest

latest: Pulling from library/python

03e1855d4f31: Already exists

a3ed95caeb02: Already exists

9269ba3950bb: Already exists

6ecee6444751: Already exists

7a0c192d4d25: Already exists

a3ed95caeb02: Already exists

66777d6149f5: Already exists

Digest: sha256:4651b83dd903ce78b1c455794f63d4108d9469a6c7fe97cd07d08a77b7e72435

Status: Image is up to date for python:latest

In [78]:

```
docker images python
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
python 689.1 MB	latest	93049cc049a6	12 days ago

In [79]:

```
docker history python
```

IMAGE SIZE	CREATED COMMENT	CREATED BY
93049cc049a6 3"]	12 days ago 0 B	/bin/sh -c #(nop) CMD ["python
<missing>	12 days ago	/bin/sh -c cd /usr/local/bin
&& ln -s easy_i	48 B	
<missing>	12 days ago	/bin/sh -c set -ex && gpg --ke
yserver ha.poo	81.53 MB	
<missing>	12 days ago	/bin/sh -c #(nop) ENV PYTHON_PI
P_VERSION=7.1.	0 B	
<missing>	12 days ago	/bin/sh -c #(nop) ENV PYTHON_VE
RSION=3.5.1	0 B	
<missing>	12 days ago	/bin/sh -c #(nop) ENV GPG_KEY=9
7FC712E4C024BB	0 B	
<missing>	12 days ago	/bin/sh -c #(nop) ENV LANG=C.UT
F-8	0 B	
<missing>	12 days ago	/bin/sh -c apt-get purge -y pyt
hon.*	978.7 kB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	314.7 MB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	122.6 MB	
<missing>	13 days ago	/bin/sh -c apt-get update && ap
t-get install	44.3 MB	
<missing>	13 days ago	/bin/sh -c #(nop) CMD ["/bin/ba
sh"]	0 B	
<missing>	13 days ago	/bin/sh -c #(nop) ADD file:e5a3
d20748c5d3dd5f	125.1 MB	

The following command will run the 'python' docker image, invoking the command 'python --version' to show us which default python version is available.

In [80]:

```
docker run python python --version
```

```
Python 3.5.1
```

Now let's look at extending this to a small 2-tier application with Flask web server as front-end and Redis as a backend database (counter) to count access to a web page.

To run this application create a Dockerfile with the following contents in the 'python_flask' directory

In [81]:

```
cd ~/src/python_flask  
cat Dockerfile
```

```
FROM python:2.7
```

```
MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>
```

```
WORKDIR /src
```

```
ADD requirements.txt /src/
```

```
RUN pip install -r requirements.txt
```

```
ADD . /src
```

```
CMD python flask_redis_app.py
```

In [82]:

```
docker build -t lab/python_flask .
```

Sending build context to Docker daemon 4.096 kB

Step 1 : FROM python:2.7

2.7: Pulling from library/python

03e1855d4f31: Already exists

a3ed95caeb02: Already exists

9269ba3950bb: Already exists

6ecee6444751: Already exists

7a0c192d4d25: Already exists

a3ed95caeb02: Already exists

7d4cc27b64c3: Pulling fs layer

Digest: sha256:84d75f33b4ae12997a0a2d51481ed18525ee7b49a5660c3301a00dfb2db19a14

Status: Downloaded newer image for python:2.7

---> 31093b2dabe2

Step 2 : MAINTAINER "Docker Build Lab" <dockerlabs@mjbright.net>

---> Running in 404474a045cb

---> 7aa58734ef0a

Removing intermediate container 404474a045cb

Step 3 : WORKDIR /src

---> Running in 1709c7d0613f

---> 50364340ba7d

Removing intermediate container 1709c7d0613f

Step 4 : ADD requirements.txt /src/

---> 2a48226d80a9

Removing intermediate container 86e962852b66

Step 5 : RUN pip install -r requirements.txt

---> Running in b6136d317c71

Collecting flask (from -r requirements.txt (line 1))

Downloading Flask-0.10.1.tar.gz (544kB)

Collecting redis (from -r requirements.txt (line 2))

Downloading redis-2.10.5-py2.py3-none-any.whl (60kB)

Collecting Werkzeug>=0.7 (from flask->-r requirements.txt (line 1))

Downloading Werkzeug-0.11.3-py2.py3-none-any.whl (305kB)

Collecting Jinja2>=2.4 (from flask->-r requirements.txt (line 1))

Downloading Jinja2-2.8-py2.py3-none-any.whl (263kB)

Collecting itsdangerous>=0.21 (from flask->-r requirements.txt (line 1))

Downloading itsdangerous-0.24.tar.gz (46kB)

Collecting MarkupSafe (from Jinja2>=2.4->flask->-r requirements.txt (line 1))

Downloading MarkupSafe-0.23.tar.gz

Building wheels for collected packages: flask, itsdangerous, MarkupSafe

Running setup.py bdist_wheel for flask

Stored in directory: /root/.cache/pip/wheels/d2/db/61/cb9b80526b8f3ba89248ec0a29d6da1bb6013681c930fca987

Running setup.py bdist_wheel for itsdangerous

Stored in directory: /root/.cache/pip/wheels/97/c0/b8/b37c320ff57e15f993ba0ac98013eee778920b4a7b3ebae3cf

Running setup.py bdist_wheel for MarkupSafe

Stored in directory: /root/.cache/pip/wheels/94/a7/79/f79a998b64c1281cb99fa9bbd33cfc9b8b5775f438218d17a7

```

Successfully built flask itsdangerous MarkupSafe
Installing collected packages: Werkzeug, MarkupSafe, Jinja2, itsdangerous, flask, redis
Successfully installed Jinja2-2.8 MarkupSafe-0.23 Werkzeug-0.11.3 flask-0.10.1 itsdangerous-0.24 redis-2.10.5
You are using pip version 7.1.2, however version 8.0.2 is available.
You should consider upgrading via the 'pip install --upgrade pip' command.
---> 7cf562861598
Removing intermediate container b6136d317c71
Step 6 : ADD . /src
---> a4044ed2f4ea
Removing intermediate container 4d4b3bce93c5
Step 7 : CMD python flask_redis_app.py
---> Running in 0d95fa7b3fa6
---> 64113b113a8a
Removing intermediate container 0d95fa7b3fa6
Successfully built 64113b113a8a

```

In [83]:

```
docker images lab/*
```

REPOSITORY	TAG	IMAGE ID	CREATED
lab/python_flask	latest	64113b113a8a	17 seconds
ago 682.9 MB			
lab/toolset	latest	6b82de537bd1	44 minutes
ago 3.245 MB			
lab/go-hello	latest	438d3567ae8a	45 minutes
ago 2.367 MB			
lab/basic	latest	e3d7946e9d93	48 minutes
ago 689.1 MB			
lab/c_prog	latest	659400ddd8d4	18 hours ago
o 877.2 kB			

Now let's run this container in the background

First we must launch our redis DB as a background container, it will be used to count the number of page accesses.

For this we launch a container instance of the 'redis' image, as a background process (-d) and we also give the instance the name 'redis' also. This will allow our 'flask' container to link to this image

In [85]:

```
docker run -d --name redis redis
```

```
b00a2d863d93bd50ea13bba2046bd6325097f05195b8d3799239a93bd5694817
```

Now we launch our 'flask' image which will be our web server.

We also launch that container in the background. Note that we link to the 'redis' container. This allows the 'flask' container to connect to the redis database to get and increment the counter value

In [86]:

```
docker run -d --link redis:db -p 5000:5000 lab/python_flask
```

f2b3fffe66eff7279d9e4d47cd15bb7e2ae6033552d5c335e6627cc8b017c44c

Now as we make successive curl (or web browser) requests to our 'flask' web server we see that page access counter gets incremented

In [87]:

```
curl http://localhost:5000
```

Hello World!
This page has been visited 1 times.
On host f2b3fffe66ef

In [88]:

```
curl http://localhost:5000
```

Hello World!
This page has been visited 2 times.
On host f2b3fffe66ef

In [89]:

```
curl http://localhost:5000
```

Hello World!
This page has been visited 3 times.
On host f2b3fffe66ef

In [90]:

```
curl http://localhost:5000
```

Hello World!
This page has been visited 4 times.
On host f2b3fffe66ef

Using Compose

[TOP](#)

Building complex systems with Compose

[TOP](#)

In [91]:

```
cd ~/src/compose
```

Create a docker-compose.yml specification file with the following contents.

This is a 2-tier architecture with

In [94]:

```
cat docker-compose.yml
```

```
version: 2
services:
  weba:
    build: ../nodeJS
    expose:
      - 80

  webb:
    build: ../nodeJS
    expose:
      - 80

  webc:
    build: ../nodeJS
    expose:
      - 80

  haproxy:
    image: haproxy
    volumes:
      - ./haproxy:/usr/local/etc/haproxy/
    links:
      - weba
      - webb
      - webc
    ports:
      - "80:80"
      - "70:70"
    expose:
      - "80"
      - "70"
```

Let's look at the docker-compose options

In [95]:

```
docker-compose
```

Define and run multi-container applications with Docker.

Usage:

```
docker-compose [-f=<arg>...] [options] [COMMAND] [ARGS...]
docker-compose -h|--help
```

Options:

-f, --file FILE	Specify an alternate compose file (default: docker-compose.yml)
-p, --project-name NAME	Specify an alternate project name (default: directory name)
--verbose	Show more output
-v, --version	Print version and exit

Commands:

build	Build or rebuild services
config	Validate and view the compose file
create	Create services
down	Stop and remove containers, networks, images, and volumes
events	Receive real time events from containers
help	Get help on a command
kill	Kill containers
logs	View output from containers
pause	Pause services
port	Print the public port for a port binding
ps	List containers
pull	Pulls service images
restart	Restart services
rm	Remove stopped containers
run	Run a one-off command
scale	Set number of containers for a service
start	Start services
stop	Stop services
unpause	Unpause services
up	Create and start containers
version	Show the Docker-Compose version information

Let's first do some cleanup of any existing containers invoked by compose

In [99]:

```
docker-compose stop
```

In [102]:

```
docker-compose rm -f
```

No stopped containers

Remove any other containers which might be listening on port 80 - here we stop and remove all

containers to be sure

In [101]:

```
docker stop $(docker ps -q)
docker rm $(docker ps -aq)
```

```
f2b3fffe66ef
b00a2d863d93
c43f44797edd
f2b3fffe66ef
b00a2d863d93
c43f44797edd
```

Now we will startup our system.

```
docker-compose up
```

will startup the containers specified in docker-compose.yml (first building any images where 'build:' is specified)

Launching in the foreground will produce the following output:

```

Creating compose_webc_1
Building weba
Step 1 : FROM node
---> baa18fdeb577
Step 2 : ADD src/ /src
---> Using cache
---> 6d6e1fb506d4
Step 3 : WORKDIR /src
---> Using cache
---> f293e728d28f
Step 4 : RUN npm install
---> Using cache
---> 16107856389c
Step 5 : EXPOSE 80
---> Using cache
---> eed510133b6b
Step 6 : CMD node index.js
---> Using cache
---> 2e6df94ecab1
Successfully built 2e6df94ecab1
Creating compose_weba_1
Creating compose_haproxy_1
Attaching to compose_webb_1, compose_webc_1, compose_weba_1, compose_haproxy_1
webb_1      | Running on http://localhost
webc_1      | Running on http://localhost
weba_1      | Running on http://localhost
haproxy_1   | [WARNING] 038/165945 (1) : Server aj_backends/weba is UP, reason: Layer7 check passed, code: 200, info: "HTTP status check returned code <3C>200<3E>", check duration: 6ms. 3 active and 0 backup servers online. 0 sessions requeued, 0 total in queue.

```

where we see that the necessary images were built and then started. We also see startup of the haproxy container and it's checks that it can contact the web servers.

Let's start the containers in the background

In [103]:

```
docker-compose up -d
```

```

Creating compose_webb_1
Creating compose_webc_1
Creating compose_weba_1
Creating compose_haproxy_1

```

Now we can see that appropriate images have been built and we can also use 'docker-compose ps' to see the running elements of our system

In [104]:

```
docker images compose*
```

REPOSITORY SIZE	TAG	IMAGE ID	CREATED
compose_weba go 645.8 MB	latest	2e6df94ecab1	4 minutes a
compose_webb go 645.8 MB	latest	2e6df94ecab1	4 minutes a
compose_webc go 645.8 MB	latest	2e6df94ecab1	4 minutes a

In [105]:

```
docker-compose ps
```

Name	Command	State	Ports
compose_haproxy_0/t10->	haproxy -f /usr/local/etc/	Up	0.0.0.0:70->7cp, 0.0.0.0:880/tcp
compose_weba_1	node index.js	Up	80/tcp
compose_webb_1	node index.js	Up	80/tcp
compose_webc_1	node index.js	Up	80/tcp

Now we can use curl connecting to localhost:80 (on which the haproxy container is listening) and we will see how haproxy is doing roundrobin scheduling to each of our web containers

In [106]:

```
curl localhost:80
```

```
<html><body>Hello from Node.js container 8727c713bfd4</body></html>
```

In [107]:

```
curl localhost:80
```

```
<html><body>Hello from Node.js container 594f34c12cb5</body></html>
```

In [108]:

```
curl localhost:80
```

```
<html><body>Hello from Node.js container 3ebf9e175082</body></html>
```

Note how the container id changes each time.

Now we can try scaling one of our services as shown below:

In [109]:

```
docker-compose scale weba=5
```

```
Creating and starting 2 ...
Creating and starting 3 ...
Creating and starting 4 ...
Creating and starting 5 ...
```

With 'docker-compose ps' we see that we now have scaled to 5 'weba' servers:

In [110]:

```
docker-compose ps
```

Name	Command	State	Ports

compose_haproxy_0/t	haproxy -f /usr/local/etc/	Up	0.0.0.0:70->70->
1	...		80/tcp
0->			cp, 0.0.0.0:80/tcp
compose_weba_1	node index.js	Up	80/tcp
compose_weba_2	node index.js	Up	80/tcp
compose_weba_3	node index.js	Up	80/tcp
compose_weba_4	node index.js	Up	80/tcp
compose_weba_5	node index.js	Up	80/tcp
compose_webb_1	node index.js	Up	80/tcp
compose_webc_1	node index.js	Up	80/tcp

In [111]:

```
curl localhost:80
curl localhost:80
curl localhost:80
curl localhost:80
curl localhost:80
curl localhost:80
```

```
<html><body>Hello from Node.js container 8727c713bfd4</body></html><html><body>Hello from Node.js container 594f34c12cb5</body></html><html><body>Hello from Node.js container 3ebf9e175082</body></html><html><body>Hello from Node.js container 8727c713bfd4</body></html><html><body>Hello from Node.js container 594f34c12cb5</body></html><html><body>Hello from Node.js container 3ebf9e175082</body></html>
```

Now let's stop our current system

In [113]:

```
docker-compose stop
docker-compose rm -f
docker-compose ps
```

No stopped containers

Name	Command	State	Ports

Building Docker

[TOP](#)

Building Docker with Docker

[TOP](#)

A major advantage of Docker is to simplify build environments.

Let's look at how we can build the Docker engine client/daemon binary without having to explicitly install a development environment.

The goal of this step is simply to show the ease with which we can build Docker, thanks to Docker itself.

We do not make particular use of the built image.

The process involves the following steps, several of which have already been performed so as to prevent excessive network utilisation during the lab. Nevertheless all steps are described here so that you can see just how easy it is to build Docker from scratch:

- Install make
- Clone the Docker source code
- Checkout the same code revision as our current Docker binary (client and daemon)
- Build the code - which pulls the docker-dev image containing the required version of the Go compiler
- Run the executable to demonstrate it is correct

Installing make

In your environment we have already installed the make package, but no compiler using yum:

```
yum install make
```

Cloning the Docker source code

We have already downloaded the Docker source code from github as follows:

```
mkdir -p /root/src/docker
cd /root/src/docker
git clone https://github.com/docker/docker .
```

To build Docker we simply have to build using the

```
make build
```

command.

Checkout the source code revision corresponding to our installed Docker Engine

If we build the latest sources this may not be compatible with our installed Docker version.

This is the case. We have 1.10.0-rc2 installed, which has API version 22, but the current github source is 1.10.0-dev which has changed to API version 23. So if we build this we find that we cannot use this client to communicate with the installed daemon.

So let's checkout the code for 1.10.0-rc2.

At the time of writing this is the latest release(candidate) of the Docker engine. We can obtain that version of the source code by referring to the releases page <https://github.com/docker/docker/releases> (<https://github.com/docker/docker/releases>) and selecting the SHA1 hash of build 1.10.0-rc2

```
git checkout c1cdc6e
```

Build the code - which pulls the docker-dev image containing the required version of the Go compiler

We can build the code as follows:

```
make build
```

We have run 'make build' already, so the docker-dev image has already been downloaded (again to prevent excessive network traffic). The docker-dev image includes the required go compiler and other build tools.

Run 'make build' again and you will see a standard build process and finally where it places the compiled binary

Run the executable to demonstrate it is correct

In preparation for the lab we built from the latest source (not the c1cdc6e version we checked out).

Run this build as follows to see that it is not compatible with the installed binary (/usr/bin/docker). We see that this binary has version 1.10.0-dev and API version 1.23 but that this cannot communicate with our installed binary which has API version 1.22.

In []:

```
cd /root/src/docker; ls -altr bundles/1.10.0-dev/binary/docker-1.10.0-dev; ./bundles/
```

But if we run our new build - as follows - created from revision c1cdc6e of the source code (corresponding to Docker version 1.10.0-rc2) we see that it has the correct version, with the same API version and can interrogate the server.

In []:

```
cd /root/src/docker; ls -altr bundles/1.10.0-rc2/binary/docker-1.10.0-rc2; ./bundles/
```

In []:

References

TOP

- [Dockerfile Reference \(https://docs.docker.com/engine/reference/builder/\)](https://docs.docker.com/engine/reference/builder/)
- [Compose file documentation \(https://docs.docker.com/compose/compose-file/\)](https://docs.docker.com/compose/compose-file/)
- [Compose file reference \(https://github.com/docker/compose/blob/1.6.0-rc1/docs/compose-file.md\)](https://github.com/docker/compose/blob/1.6.0-rc1/docs/compose-file.md)
- [Visualizing Docker Containers and Images \(http://merrigrove.blogspot.in/2015/10/visualizing-docker-containers-and-images.html\)](http://merrigrove.blogspot.in/2015/10/visualizing-docker-containers-and-images.html)
- [Awesome Docker \(https://github.com/veggie Monk/awesome-docker\)](https://github.com/veggie Monk/awesome-docker)
- [Docker Cheat Sheet \(\)](#)
- [Building Good Docker Images \(http://jonathan.bergknoff.com/journal/building-good-docker-images\)](http://jonathan.bergknoff.com/journal/building-good-docker-images)
- [How to scale a Docker Container with Docker Compose \(https://www.brianchristner.io/how-to-scale-a-docker-container-with-docker-compose/\)](https://www.brianchristner.io/how-to-scale-a-docker-container-with-docker-compose/)
- [Docker Compose Demo \(https://github.com/vegasbrianc/docker-compose-demo\)](https://github.com/vegasbrianc/docker-compose-demo)

In []: